

YASHAS MAHANTESH PAREESHWAD

Belagavi, Karnataka, India



+91 9538876675



yashas422005@gmail.com

LinkedIn

GitHub

Summary

AI/ML Engineer specializing in building end-to-end automation pipelines, LLM-powered systems, and scalable backend architectures. Experienced in retrieval-augmented generation (RAG), AI agents, vector databases, and real-time inference systems. Strong foundation in n8n, LangChain, and LangGraph, with hands-on experience deploying production-grade AI systems on a VPS.

Technical Skills

Languages: Python

Databases: PostgreSQL, Redis, Supabase (pgvector)

AI/ML: Machine Learning, Deep Learning, NLP, Generative AI, Agentic AI

LLM Systems: LLMs, RAG, AI Agents, Prompt Engineering

Frameworks & Tools: n8n, LangChain, LangGraph, TensorFlow, Scikit-learn, Docker

Experience

Vijesha IT Services

6 Months

Intern

- Worked on a project titled “Sentiment Analysis using DistilBERT and VADER,” comparing rule-based and transformer-based approaches to sentiment classification.
- Implemented and benchmarked VADER (lexicon-based) against DistilBERT (transformer-based) models, evaluating differences in accuracy, contextual understanding, and scalability.
- Documented the evolution of sentiment analysis techniques from lexicon-based methods to deep learning, and presented findings on how modern NLP models outperform traditional approaches on complex language patterns.

Projects

Electronics Product Q&A Agent (Agentic RAG) | n8n, Google Gemini, Supabase (pgvector), Google Sheets

- Built a conversational Agentic RAG system enabling natural-language product queries for electronics including mice, keyboards, and webcams, retrieving relevant information from a vector store.
- Designed an automated ingestion pipeline that extracts product catalog data from Google Sheets, chunks and embeds product records using Gemini embeddings, and indexes them in Supabase (pgvector) with category and SKU metadata for filtered retrieval.
- Implemented semantic query caching to match incoming questions by meaning rather than exact text, reducing redundant LLM and retrieval calls and **improving average response time by 36%**.

Agentic HR Recruitment Assistant | n8n, Google Gemini, Supabase (pgvector), PostgreSQL, Redis, Gmail API, Google Calendar API

- Designed a multi-agent recruitment system where a central conversational agent delegates tasks to specialized agents for candidate retrieval, email drafting, and interview scheduling, **reducing manual resume review from ~20–25 minutes to seconds per candidate**.
- Built a RAG pipeline over candidate resumes using Gemini embeddings and Supabase (pgvector), enabling recruiters to search and shortlist candidates using natural-language queries.
- Automated interview scheduling through a calendar agent that checks real-time Google Calendar availability and books meetings directly from the recruiter chat.
- Implemented Redis-based semantic response caching and conversational memory to reduce redundant LLM calls and maintain context across recruitment sessions.
- Automated personalized shortlisting and interview emails using a Gmail-integrated agent that drafts candidate communications and includes generated interview links.

Education

B.E. Computer Science

2027

KLS Gogte Institute of Technology, Belagavi

CGPA: 9(current) / 10

PUC (12th Grade)

2023

Raja Lakhamagowda Science PU College

88%